

BRANDON HAN

AI / ML Engineer | LLM Evaluation, AI Agents, Customer Support Automation
https://www.thehanbrand.dev | (626)-404-4082 | brandonh4n@gmail.com | [LinkedIn](#) | [Github](#)

EDUCATION

University of Southern California Los Angeles, CA
Master of Science, Computer Science January 2024-May 2026
Relevant Courses: Advanced Databases | Algorithms | Machine Learning | Deep Learning & Optimization | Web Technologies

University of California, Los Angeles Los Angeles, CA
Bachelor of Science, Biochemistry January 2017-December 2019

TECHNICAL SKILLS

Programming Languages: Python | Java | C++ | C | JavaScript | TypeScript | SQL | Go
ML and NLP: PyTorch | TensorFlow | Hugging Face Transformers | RoBERTa | scikit-learn
Data and Evaluation: Pandas | NumPy | VADER | Model Evaluation
Engineering: Flask | FastAPI | Node.js | React | PostgreSQL | MySQL | MongoDB | Docker | AWS | Git

PROJECTS

Bias Induced News Generation with LLMs | PyTorch, Hugging Face, RoBERTa Los Angeles, CA
Machine Learning Engineer May 2026

- Engineered an automated evaluation pipeline to expose and quantify political bias across multiple LLM architectures
- Built a pipeline for the ingestion, cleaning, and labeling of multi-partisan news snippets for LLM evaluation
- Processed 14,000 biased articles to automate the generation and scoring 1,700 test cases

Stock Sentiment Tracker | Python, Hugging Face, Flask, SHAP Los Angeles, CA
Machine Learning Engineer November 2025

- Fine-tuned FinBERT for informal finance text by relabeling StockEmotions with Twitter-RoBERTa and training on social-market language
- Improved accuracy on informal finance text by 57% against baseline setup, then benchmarked using MSE, MAE
- Implemented SHAP to inspect model predictions and identify finance-specific sentiment cues that were easy to miss in standard language models

Weenix Operating System Kernel Development | C, x86 Assembly, GNU Make, QEMU Los Angeles, CA
Software Engineer February 2025

- Engineered foundational system components for a 32-bit architecture
- Developed process and thread life-cycle management, context switching, thread bootstrap, and scheduler queue primitives
- Integrated software with emulated hardware (QEMU) using C and Assembly, navigating complex system constraints, physical memory allocation, and concurrency.

EXPERIENCE

Wasabi Cloud Technologies Boston, MA
AI Engineering Intern (Remote) June 2025-August 2025

- Built a React escalation dashboard that combined customer requests, upload failures, account activity, and deployment notes into one timeline, reducing median support-to-engineering handoff time by 23%
- Created labeled datasets from historical support tickets, customer notes, and internal business text to improve AI triage accuracy on company-specific language and recurring cloud-storage issues
- Built an evaluation loop for AI recommendations using resolved incidents, tracking false positives, missing context, and low-confidence cases that required human review

University of California, Los Angeles Los Angeles, CA
Research Operations Manager January 2019-December 2023

- Led rollout of GPU-backed compute and multi-vendor lab systems, translating wet-lab constraints into implementation plans, integration tests, and researcher onboarding
- Secured \$1.2M in infrastructure funding for automation, imaging, and compute platforms through technical proposals and systems planning
- Contributed to 5 peer-reviewed publications involving ML implementation, imaging pipelines, and research infrastructure