

BRANDON HAN

ML / AI Engineer | Applied NLP, Computer Vision, Production ML Systems

<https://www.thehanbrand.dev> | (626)-404-4082 | brandonh4n@gmail.com | [LinkedIn](#) | [Github](#)

EDUCATION

University of Southern California

Los Angeles, CA

Master of Science

January 2024-May 2026

Relevant Courses: Advanced Databases | Algorithms | Machine Learning | Deep Learning & Optimization | Web Technologies

University of California, Los Angeles

Los Angeles, CA

Bachelor of Science

January 2017-December 2019

TECHNICAL SKILLS

Programming Languages: Python | Java | C++ | C | JavaScript | TypeScript | SQL | Go

ML and NLP: PyTorch | TensorFlow | Hugging Face Transformers | RoBERTa | scikit-learn

Data and Evaluation: Pandas | NumPy | VADER | Model Evaluation

Engineering: Flask | FastAPI | Node.js | React | PostgreSQL | MySQL | MongoDB | Docker | AWS | Git

PROJECTS

Bias Induced News Generation with LLMs | PyTorch, Hugging Face, RoBERTa

Los Angeles, CA

Machine Learning Engineer

May 2026

- Engineered an automated evaluation pipeline to expose and quantify political bias across multiple LLM architectures
- Built a pipeline for the ingestion, cleaning, and labeling of multi-partisan news snippets for LLM evaluation
- Processed 14,000 biased articles to automate the generation and scoring of 1,700 test cases

Stock Sentiment Tracker | Python, Hugging Face, Flask, SHAP

Los Angeles, CA

Machine Learning Engineer

November 2025

- Fine-tuned FinBERT for informal finance text by relabeling StockEmotions with Twitter-RoBERTa and training on social-market language
- Improved accuracy on informal finance text by 57% against baseline setup, then benchmarked using MSE, MAE
- Implemented SHAP to inspect model predictions and identify finance-specific sentiment cues that were easy to miss in standard language models

Weenix Operating System Kernel Development | C, x86 Assembly, GNU Make, QEMU

Los Angeles, CA

Software Engineer

February 2025

- Engineered foundational system components for a 32-bit architecture
- Developed process and thread life-cycle management, context switching, thread bootstrap, and scheduler queue primitives
- Integrated software with emulated hardware (QEMU) using C and Assembly, navigating complex system constraints, physical memory allocation, and concurrency.

EXPERIENCE

Wasabi Cloud Technologies

Boston, MA

AI Engineering Intern (Remote)

June 2025-August 2025

- Built a React dashboard that brought customer requests, upload failures, and deploy notes into one escalation timeline, reducing median incident handoff time by 23%
- Created labeled datasets from historical support tickets, customer notes, and internal business text to improve AI model accuracy on company-specific language and recurring issue patterns
- Developed a first-pass AI triage workflow that converted ticket text and error patterns into plain-language diagnostic notes for support teams

University of California, Los Angeles

Los Angeles, CA

Research Operations Manager

January 2019-December 2023

- Led rollout of GPU-backed compute and multi-vendor lab systems, translating wet-lab constraints into implementation plans, integration tests, and researcher onboarding
- Secured \$1.2M in infrastructure funding for automation, imaging, and compute platforms through technical proposals and systems planning
- Contributed to 5 peer-reviewed publications involving ML implementation, imaging pipelines, and research infrastructure